

AI USED: CHAT GPT

PERCENTAGE: 20%

SCREENSHOTS ARE BELOW:

I DONT REALLY UNDERSTAND HOW TO START THIS.....MAYBE EXPLAINING TO DETAILS CAN GUIDE ME
Each row is a completed taxi trip. Your goal is to predict the tip a passenger leaves (tip_amount, a continuous value) from the trip's characteristics. The available columns are:

VendorID, passenger_count, trip_distance, RatecodeID, store_and_fwd_flag, PULocationID, DOLocationID, payment_type, fare_amount, mta_tax, tolls_amount, improvement_surcharge, tip_amount

Show more ▾

No problem — this is actually the **most important starting point of the whole lab**, so I'll make it very clear.

You are basically doing **3 things only** in Part 1.1:

1. "Load the data"
2. "Understand what is inside"
3. "Look at the target (tip_amount)"

BIG IDEA (what you are doing)

You are asking:

"What does this taxi dataset look like before I build my model?"

So you are NOT modeling yet, just exploring

+ Ask anything

 Snipping Tool

Screenshot copied
Automatically saved

EXPLAIN TO ME WHAT I AM SUPPOSED TO DO PLS

TODO: Handle missing / invalid rows

e.g. drop or impute NaNs; remove rows where trip_distance <= 0 or fare_amount <= 0;

decide how to treat negative tip_amount values.

TODO: Feature engineering — create at least ONE new, justified feature, for example:

- tip_rate proxy, fare-per-mile = fare_amount / trip_distance

- total surcharges = mta_tax + tolls_amount + improvement_surcharge

Show more ▾

Good — this is the part where most people get confused. I'll break it down into **very simple actions** so you know exactly what to do.

Think of this step as:

“Clean the taxi data so a machine learning model can understand it.”

🧠 PART 1 — Handle missing + invalid data

? What are you doing?

You are removing “bad rows” that don't make sense.

🚫 Examples of bad data:

- missing values (NaN)
- trip_distance ≤ 0 (impossible trip)



+ Ask anything



EXPLAIN THE CODE VERY WELLLAYMAN TERMS

Alright — I'll explain it in **very simple, everyday language**, like you've never seen machine learning before.

Think of this whole section as:

🔪 "Cleaning and preparing taxi data so a computer can learn from it."

🧠 1. Handling missing + invalid data

Code:

Python

```
taxi = taxi.dropna()
taxi = taxi[taxi["trip_distance"] > 0]
taxi = taxi[taxi["fare_amount"] > 0]
taxi = taxi[taxi["tip_amount"] >= 0]
```



Run

💡 What this REALLY means:

Imagine your taxi data is like a notebook filled with trip records.

Some pages are:

- blank (missing data)
- nonsense (trip distance = 0 → no trip happened)

...

+ Ask anything



ChatGPT can make mistakes. Check important info.

FINAL SPLIT

You now have:

- 60% → Train
- 20% → Validation
- 20% → Test

STEP 4 — Scaling (VERY IMPORTANT RULE)

 Python



 Run

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train) # Learn from train ONLY
X_val_scaled = scaler.transform(X_val)
X_test_scaled = scaler.transform(X_test)
```

WHY “FIT ONLY ON TRAIN”?

Simple meaning:

You must not let the model “see” validation or test data during learning.

Otherwise it's like:

 studying with exam answers already visible



 Ask anything



WHAT IS IT ABOUT

I WANT TO KNOW AND UNDERSTAND WHAT TO DO HERE

TODO: Use ONLY the scaled obesity features (no target) for clustering.

TODO: Choose the number of clusters k with the Elbow method (and/or silhouette score).

from sklearn.cluster import KMeans

from sklearn.metrics import silhouette_score

Loop k = 2..10, record inertia_ (and silhouette), and plot inertia vs k.

TODO: Fit KMeans with your chosen k

Show more ▾

Good — this is the **core idea of your unsupervised learning section**, so I'll explain it in plain, simple terms first.

WHAT THIS SECTION IS ABOUT (BIG IDEA)

You are pretending:

"I DO NOT KNOW the obesity labels (NObeyesdad)."

Then you ask:

"Can the computer still find natural groups of people just from their features?"

That's it.

This is called:



+ Ask anything



I DID SOMETHING BUT IT IS NOT WORKING AND IT DOESNT REALLY MAKE SENSE....COULD U GIVE ME YOURS SO I CAN RECTIFY MY MISTAKE

ot you — I'll give you a **clean working version** of the whole K-Means section so you can compare and fix ours.

Il also keep it simple so you see exactly what each part is doing.

✓ FULL WORKING K-MEANS PIPELINE (COPY THIS)

. Use **ONLY** features (remove target)

Python



Run

```
X_cluster = obesity_clean.drop("NObeyesdad", axis=1)
```

. Find best k (Elbow + Silhouette)

Python



Run

```
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
```



+ Ask anything



ChatGPT can make mistakes. Check important info.

YOUR OWN IS VERY DIFFERENT FROM MINE.....EXPLAIN IT CODE BY CODE

Good — let's slow this right down and go **line by line exactly what each piece of code is doing**, in plain language. Forget "ML theory" for now — think of it like following a cooking recipe.

I'll also show you what each line *means in real life*.

STEP 1 — Remove the answer column

 Python



 Run

```
X_cluster = obesity_clean.drop("NObesidad", axis=1)
```

What this means:

You are saying:

"I want ONLY the input information, not the answers."

Think of it like:

You give a student data:

...

- Age

+ Ask anything



ChatGPT can make mistakes. Check important info.

🧩 STEP 5 — Build KMeans model

Python



Run

```
kmeans = KMeans(n_clusters=k, random_state=RANDOM_STATE, n_init=10)
```

💡 Meaning:

You are saying:

"Split data into k groups"

Break it down:

- `n_clusters=k` → number of groups
- `random_state` → ensures same results every time
- `n_init=10` → tries multiple starts for better grouping

Think of it like:

You are asking:

"Split these people into k groups and try 10 different ways to do it, then pick the best one."

🧩 STEP 6 — Assign clusters

Python



Run

+ Ask anything

